

07. Diffusion Theory

Overview

Diffusion models define a forward noising process and learn to reverse it. This module derives the variance schedule, closed-form noisy samples, and the connection between score matching and denoising objectives. You need this foundation before DDPM sampling and latent diffusion in the next module.

Lab: Lab 07.

Prior modules: Module 05.

Contents

1	From GANs to Diffusion	2
2	Forward Diffusion Process	2
3	Reverse Process Goal	2
4	Variational Lower Bound	2
5	Noise Schedules	2
6	Score Matching Connection	2
7	Lab	3
8	Signal-to-Noise Ratio	3
8.1	Stochastic vs. Deterministic Sampling	3
9	Extended Intuition	3
9.1	Score Matching Connection	4
10	Numerical Walkthrough	4
11	PyTorch Implementation Notes	4
12	Local GPU Constraints	5
13	Debugging Checklist	5
14	Practice Problems	5

Module track

This module builds on **Module 05**. Read those PDFs first. References to other modules appear as **Module NN** (not page numbers, because each module is its own PDF).




Figure 1: Module sequence (solid box: this PDF; dashed: typical next step).

1 From GANs to Diffusion

While GANs (**Module 05** (GAN objective)) learn implicit samplers, **diffusion models** define a forward noising process and learn its reverse. Foundations support DDPM [2] in **Module 08** (DDPM loss).

2 Forward Diffusion Process

Define Markov chain adding Gaussian noise over $t = 1, \dots, T$:

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I}), \quad (1)$$

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}), \quad (2)$$

where $\alpha_t = 1 - \beta_t$, $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$.

Note. The closed form $q(\mathbf{x}_t | \mathbf{x}_0)$ enables single-step noising during training, no sequential simulation required.

3 Reverse Process Goal

We learn $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$ to denoise. For small β_t , reverse is also Gaussian:

$$p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \boldsymbol{\mu}_\theta(\mathbf{x}_t, t), \boldsymbol{\Sigma}_\theta(\mathbf{x}_t, t)). \quad (3)$$

4 Variational Lower Bound

Maximize ELBO on $\mathbf{x}_0$:

$$\log p_\theta(\mathbf{x}_0) \geq \mathbb{E}_q \left[\log \frac{p_\theta(\mathbf{x}_{0:T})}{q(\mathbf{x}_{1:T} | \mathbf{x}_0)} \right] = \mathcal{L}_{\text{VLB}}. \quad (4)$$

Decomposes into reconstruction, prior matching, and denoising terms across timesteps.

5 Noise Schedules

Linear: β_t from 10^{-4} to 0.02 over $T = 1000$. Cosine schedule improves low-noise regions:

$$\bar{\alpha}_t = \frac{f(t)}{f(0)}, \quad f(t) = \cos^2 \left(\frac{t/T + s}{1 + s} \cdot \frac{\pi}{2} \right). \quad (5)$$


6 Score Matching Connection

The score $\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t)$ relates to predicted noise $\boldsymbol{\epsilon}_\theta$:

$$\nabla_{\mathbf{x}_t} \log q(\mathbf{x}_t | \mathbf{x}_0) = -\frac{\boldsymbol{\epsilon}}{\sqrt{1 - \bar{\alpha}_t}}, \quad \mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}. \quad (6)$$

Table 1: Noise schedule comparison.

Schedule	Simplicity	Sample quality
Linear	High	Good
Cosine	Medium	Better (low SNR)
Learned	Low	Active research

**Figure 2:** Forward noising (solid) follows $q(\mathbf{x}_t | \mathbf{x}_{t-1})$; training learns the reverse kernel p_θ (dashed), as in DDPM [2].

7 Lab

Lab 07 visualizes $q(\mathbf{x}_t | \mathbf{x}_0)$ from (1) across t . Implement DDPM training objective in **Module 08** (DDPM loss).

8 Signal-to-Noise Ratio

Define $\text{SNR}(t) = \bar{\alpha}_t / (1 - \bar{\alpha}_t)$; denoising difficulty increases with t . Predicting $\mathbf{x}_0$ directly relates to noise via Tweedie’s formula:

$$\mathbb{E}[\mathbf{x}_0 | \mathbf{x}_t] = \frac{1}{\sqrt{\bar{\alpha}_t}} \left(\mathbf{x}_t - \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon} \right). \quad (7)$$

8.1 Stochastic vs. Deterministic Sampling

DDIM skips substeps using non-Markovian inference for faster generation.

9 Extended Intuition

Diffusion models define a forward process that gradually adds Gaussian noise until data becomes nearly pure noise. The reverse process learns to denoise step by step, recovering structure from $\mathbf{x}_T \approx \mathcal{N}(\mathbf{0}, \mathbf{I})$. Instead of adversarial balance, training reduces to supervised noise prediction (or score matching) at random timesteps.

The variance schedule β_t controls how much noise enters per step. Small β_t over many steps ($T = 1000$) yields a smooth corruption path. Ho et al. use linear schedule from $\beta_1 = 10^{-4}$ to $\beta_T = 0.02$; cosine schedules often improve low-noise steps.

Closed-form sampling of $\mathbf{x}_t$ given $\mathbf{x}_0$ avoids iterating the forward chain during training:

$$q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N}\left(\sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I}\right), \quad (8)$$

with $\alpha_t = 1 - \beta_t$ and $\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$. The network typically predicts noise $\boldsymbol{\epsilon}$ that was used to construct $\mathbf{x}_t$, not $\mathbf{x}_0$ directly. DDPM in **Module 08** turns this into a trainable sampler with discrete reverse updates.

9.1 Score Matching Connection

The score function $\nabla_{\mathbf{x}} \log p(\mathbf{x})$ points toward higher density regions. Denoising score matching trains a network to predict noise ϵ which is proportional to the score under Gaussian corruption. Signal-to-noise ratio $\text{SNR}(t) = \bar{\alpha}_t / (1 - \bar{\alpha}_t)$ drops exponentially in t for linear schedules; low-SNR steps are hardest to learn. Training uniformly over t oversamples easy high-noise steps relative to perceptual importance; some schedules reweight loss by SNR or use cosine β_t to balance. Conditional diffusion adds class or text embedding to ϵ_θ input; classifier-free guidance in **Module 08** builds on the same dropout trick at train time. Forward process is Markov: $q(\mathbf{x}_t | \mathbf{x}_{t-1})$ depends only on the previous step, enabling single-step training samples at any t .

10 Numerical Walkthrough

1D toy for intuition: $\mathbf{x}_0 = [2.0]$, $\beta_t = 0.1$, so $\alpha_t = 0.9$, $\bar{\alpha}_1 = 0.9$.

Noisy sample: $\mathbf{x}_1 = \sqrt{0.9} \cdot 2.0 + \sqrt{0.1} \epsilon = 1.897 + 0.316\epsilon$. If $\epsilon = 0.5$, $\mathbf{x}_1 \approx 2.05$ (still close to signal).

After $t = 10$ with constant $\beta = 0.1$: $\bar{\alpha}_{10} = 0.9^{10} \approx 0.349$, signal scale $\sqrt{\bar{\alpha}_{10}} \approx 0.59$. After $t = 50$: $\bar{\alpha}_{50} \approx 0.005$, mean coefficient on $\mathbf{x}_0$ is $\sqrt{0.005} \approx 0.07$: mostly noise.

Training draw: pick $t \sim \text{Uniform}\{1, \dots, T\}$, sample $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$, build $\mathbf{x}_t$, regress $\|\epsilon - \epsilon_\theta(\mathbf{x}_t, t)\|^2$. For MNIST $\mathbf{x}_0 \in \mathbb{R}^{784}$, $T = 1000$, $B = 64$, a small MLP noise predictor might need 200k steps to produce blurry digits before conv models help.

At $t = 500$ with linear schedule to $\beta_T = 0.02$: $\bar{\alpha}_{500} \approx \exp(-\sum \beta) \ll 1$; the network must infer structure from mostly noise. Predicting $\mathbf{x}_0$ directly is equivalent algebraically to predicting ϵ , but ϵ targets have unit variance at every t , which stabilizes MSE loss scaling. Always sample a fresh $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ each training step; reusing noise couples batches incorrectly.

Forward diffusion visualization: plot $\mathbf{x}_t$ at $t \in \{0, 100, 500, 999\}$ for one MNIST digit; human eye loses structure around t where $\sqrt{\bar{\alpha}_t} < 0.3$. Reverse sampling is not trained explicitly; quality depends entirely on whether ϵ_θ generalizes across all t , especially mid-range SNR.

11 PyTorch Implementation Notes

- Precompute $\bar{\alpha}_t$, $\sqrt{\bar{\alpha}_t}$, $\sqrt{1 - \bar{\alpha}_t}$ buffers; register with `register_buffer` so they move with device.
- Timestep t embedding: sinusoidal like Transformers, or `nn.Embedding(T, dim)` indexed by integer t .
- Training noise: `x_t = sqrt_ab[t]*x0 + sqrt_1m_ab[t]*eps` with broadcast shapes `(B,1,1,1)` for images.
- Loss: `F.mse_loss(eps, model(x_t, t))`; optional weighting by signal-to-noise ratio at advanced stage.
- Do not backprop through the forward diffusion sampling loop during training; single-step prediction only.

```
t = torch.randint(0, T, (B,), device=x0.device)
eps = torch.randn_like(x0)
x_t = extract(sqrt_alphas_cumprod, t, x0.shape) * x0 \
      + extract(sqrt_one_minus_alphas_cumprod, t, x0.shape) * eps
loss = F.mse_loss(model(x_t, t), eps)
```

12 Local GPU Constraints

Theory module math runs on CPU; preview training a tiny MLP denoiser on flattened MNIST fits in 4 GB easily. Store schedule tensors on GPU once; avoid recomputing $\bar{\alpha}_t$ each batch.

When jumping to conv U-Net in **Module 08**, 28×28 , $B = 128$ is feasible; 64×64 needs $B \leq 32$ on a 1650. Mixed precision safe for MSE noise prediction in most setups.

Cosine schedule (Nichol & Dhariwal): $\bar{\alpha}_t = \cos^2\left(\frac{t/T+s}{1+s} \cdot \frac{\pi}{2}\right)$ with small s prevents β_t from blowing up at $t = 0$. ELBO connection: DDPM training approximates variational bound on log-likelihood; simplified MSE drops weighting terms that matter at low t but suffice for MNIST labs. Score-based SDE view: forward SDE adds noise continuously; reverse SDE uses learned score $\nabla_{\mathbf{x}} \log p_t(\mathbf{x})$; DDPM is a discretization (continuous theory in **Module 08** sampling). Variance-preserving (VP) vs variance-exploding (VE) schedules change how $\bar{\alpha}_t$ behaves at large t ; Ho DDPM is VP; know which your code implements.

13 Debugging Checklist

- Loss decreases but samples are noise: reverse sampler not implemented or wrong $\bar{\alpha}_t$ indexing off-by-one.
- Exploding $\mathbf{x}_t$ values: forgot to normalize images to $[-1, 1]$ or unit variance before noising.
- Timestep embedding ignored: model output independent of t ; verify embedding concatenation or AdaGN injection.
- Schedule too aggressive: β_T too large makes $\bar{\alpha}_T \approx 0$ too early; try cosine schedule.
- Shape bug on broadcast: `t` shape $(B,)$ must expand to image dims via helper `extract`.
- Loss flat at 1.0: model ignoring t embedding; verify sinusoidal or learned time index is concatenated into hidden state.

14 Practice Problems

Problem 1. With $\beta_t = 0.02$ constant, compute $\bar{\alpha}_{100}$ and interpret signal fraction $\sqrt{\bar{\alpha}_{100}}$.

Problem 2. Derive $\mathbf{x}_t$ as linear combo of $\mathbf{x}_0$ and ϵ given the Gaussian forward kernel.

Problem 3. Plot $\sqrt{\bar{\alpha}_t}$ vs t for linear vs cosine schedules on the same plot.

Problem 4. Train a scalar denoiser on 1D mixture of Gaussians and visualize reverse samples at $t = 100, 500, 999$.

Problem 5. For fixed $\mathbf{x}_0$ and ϵ , plot $\|\mathbf{x}_t - \mathbf{x}_0\|_2$ vs t . Compare linear vs cosine $\bar{\alpha}_t$ schedules.

Problem 6. Implement cosine β_t schedule and $\log \bar{\alpha}_T$. Confirm $\bar{\alpha}_T > 0$ so $q(\mathbf{x}_T | \mathbf{x}_0)$ retains slight signal scale.

Denoising at each t is band-pass filtering: high t removes all frequencies; the network restores structure frequency-by-frequency across training steps. Hybrid loss weighting $\lambda(t)$ by SNR is optional; start with uniform t sampling before tuning per-timestep weights.

Visualize forward noise in **Lab 07**; full DDPM loop in **Module 08**.

Source certification

This module lists where each core definition, equation, figure, and summary table comes from. Pedagogical simplifications (lab batch sizes, epoch counts, illustrative plots) are labeled in extension sections and are not claimed as optimal production settings.

Table 3: Certified topic map for Module 07.

Topic	Primary source	Where in this PDF
Forward diffusion $q(\mathbf{x}_t \mid \mathbf{x}_{t-1})$	[2]	Eq. 1
Closed-form $q(\mathbf{x}_t \mid \mathbf{x}_0)$	[2]	Eq. 1
Reverse kernel p_θ	[2]	Section 3
Linear / cosine schedules	[3]	Section 5
Score matching link	[5]	Section 6
Diffusion chain figure	[2]	Fig. 2

Full bibliographic entries appear below. Figures are schematic unless a caption cites a specific paper figure. If a statement is not listed here, treat it as general ML practice or lab-specific guidance, not a cited theorem.

References

- [1] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, et al. Generative adversarial nets. In *NeurIPS*, 2014.
- [2] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *NeurIPS*, 2020.
- [3] Alexander Quinn Nichol and Prafulla Dhariwal. Improved denoising diffusion probabilistic models. In *ICML*, 2021.
- [4] Jascha Sohl-Dickstein, Eric Weiss, Niru Maheswaranathan, and Surya Ganguli. Deep unsupervised learning using nonequilibrium thermodynamics. In *ICML*, 2015.
- [5] Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, et al. Score-based generative modeling through stochastic differential equations. In *ICLR*, 2021.