

09. BERT

Overview

BERT pre-trains bidirectional context with masked language modeling and next-sentence prediction, then fine-tunes on downstream tasks. This module explains tokenization with WordPiece, segment embeddings, the [CLS] pooling head, and fine-tuning hyperparameters for classification.

Lab: Lab 09.

Prior modules: Module 04.

Contents

1	Pre-training Paradigm	2
2	Model Architecture	2
3	Masked Language Modeling	2
4	Next Sentence Prediction	2
5	Fine-Tuning Recipe	2
6	Comparison with Seq2seq Encoders	2
7	Lab	3
8	Tokenization and Subwords	3
8.1	Whole-Word Masking	3
9	Extended Intuition	3
9.1	MLM Mechanics Beyond the Paper Diagram	4
10	Numerical Walkthrough	4
11	PyTorch Implementation Notes	4
12	Local GPU Constraints	5
13	Debugging Checklist	5
14	Practice Problems	5

Module track

This module builds on **Module 04**. Read those PDFs first. References to other modules appear as **Module NN** (not page numbers, because each module is its own PDF).



Figure 1: Module sequence (solid box: this PDF; dashed: typical next step).

1 Pre-training Paradigm

BERT [1] pre-trains deep bidirectional representations from unlabeled text, then fine-tunes on downstream tasks. Uses Transformer encoder blocks from **Module 04** (transformer block).

2 Model Architecture

BERT-base: $L = 12$, $d = 768$, $h = 12$, ≈ 110 M parameters. Input representation sums token, segment, and position embeddings:

$$\mathbf{E}_{\text{input}} = \mathbf{E}_{\text{token}} + \mathbf{E}_{\text{segment}} + \mathbf{E}_{\text{position}}. \quad (1)$$

[CLS] token embedding aggregates sequence-level predictions.

Bidirectional context. Unlike causal LMs in **Module 10** (causal LM objective), BERT attends to full left and right context via masked language modeling.

3 Masked Language Modeling

Randomly mask 15% of tokens; of those, 80% replaced with [MASK], 10% random token, 10% unchanged.

$$\mathcal{L}_{\text{MLM}} = - \sum_{i \in \mathcal{M}} \log P(x_i | \tilde{\mathbf{x}}), \quad (2)$$

where $\mathcal{M}$ is the set of masked positions.

4 Next Sentence Prediction

Binary classification: is sentence B the actual next sentence after A ?

$$P(\text{IsNext} | A, B) = \text{softmax}(\mathbf{W}_{\text{NSP}} \mathbf{h}_{[\text{CLS}]}). \quad (3)$$

Later work (RoBERTa) shows NSP is optional for many tasks.

5 Fine-Tuning Recipe

Add task head on [CLS] (classification) or token outputs (NER). Hyperparameters: lr 2×10^{-5} – 5×10^{-5} , batch 16 to 32, 2 to 4 epochs.


6 Comparison with Seq2seq Encoders

BERT encoder parallels bidirectional LSTM encoders in **Module 02** (bidirectional LSTM) but with full self-attention in **Module 04** (self-attention).

$$\text{F1}_{\text{QA}} = \frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}. \quad (4)$$

Table 1: BERT fine-tuning heads.

Task	Head
Classification	Linear on [CLS]
Token labeling	Linear per token
QA	Span start/end logits

**Figure 2:** BERT input assembly.

7 Lab

Lab 09 fine-tunes BERT on sentiment; compare to GPT-style fine-tuning in **Module 10** (causal LM objective). Parameter-efficient tuning via LoRA in **Module 11** (LoRA definition).

8 Tokenization and Subwords

WordPiece/BPE splits rare words; vocabulary $\mathcal{V}$ typically 30k tokens. Max sequence length 512 limits quadratic attention cost from **Module 04**.

8.1 Whole-Word Masking

Mask entire word pieces belonging to one word, improves MLM coherence vs. random subword masks.

9 Extended Intuition

BERT pre-trains on masked language modeling (MLM): randomly mask 15% of tokens, predict them using bidirectional context. Unlike left-to-right LMs, BERT sees tokens to the left and right of a mask simultaneously through self-attention without causal masking. Next sentence prediction (NSP) was in the original paper; many later checkpoints drop it, but the idea was to learn document-level coherence.

Fine-tuning adds a task head on top of the [CLS] token representation for classification, or token heads for tagging. WordPiece tokenization splits rare words into subwords; the vocabulary (30k for **bert-base**) balances coverage and table size. Learning rates stay small (2×10^{-5} to 5×10^{-5}) because pretrained weights already encode syntax and semantics.

Compare with GPT-style causal models in **Module 10**: BERT excels at understanding tasks (classification, NER), GPT at generation. LoRA in **Module 11** reduces memory when fine-tuning large BERT variants on consumer GPUs.

9.1 MLM Mechanics Beyond the Paper Diagram

Of the 15% selected tokens, BERT replaces 80% with [MASK], 10% with random tokens, 10% unchanged. That trick forces representations that do not assume every input id is always visible at train time. Whole-word masking (later checkpoints) masks all WordPiece pieces of a word together, aligning MLM targets with linguistic units. At fine-tune time you never feed [MASK] tokens; only pre-training uses them, so do not confuse MLM with fill-in-the-blank inference. ALBERT uses factorized embeddings and cross-layer parameter sharing; smaller footprint than `bert-base` when VRAM blocks full fine-tune before **Module 11**. NSP (next sentence prediction) pre-training is deprecated in RoBERTa-style pipelines; MLM alone drives most downstream gains on GLUE-scale tasks.

10 Numerical Walkthrough

`bert-base-uncased`: $L = 12$, $d = 768$, heads $H = 12$, FFN 3072, vocab $V = 30522$, max length $T = 128$, $B = 16$.

Token ids shape $(16, 128)$; embeddings sum token, segment, and position vectors $\in \mathbb{R}^{768}$. One self-attention layer: Q/K/V projections $768 \rightarrow 768$ each; attention $O(T^2 d) = 128^2 \cdot 768 \approx 12.6\text{M}$ multiply-adds per layer per sequence (rough order).

Fine-tune SST-2 binary sentiment: [CLS] vector $\mathbf{h}_{\text{CLS}} \in \mathbb{R}^{768}$ passed through `Linear(768, 2)`. Cross-entropy loss on batch; 2094 train examples, 2 epochs, $\text{lr} = 3 \times 10^{-5}$, warmup 10% steps. Expect val accuracy 90 to 93% with proper max-length padding and attention mask.

Trainable params full fine-tune: $\approx 110\text{M}$ (entire BERT + head). AdamW memory $\approx 3 \times$ params in fp32 $\rightarrow 1.3\text{GB}+$ weights alone; activations push total past 4 GB at $B = 16$, $T = 128$.

Segment embeddings (`token_type_ids`) distinguish sentence A vs B in pair tasks; for single-sentence SST-2, set all zeros. Warmup-linear decay: first 10% of steps ramp lr from 0 to 3×10^{-5} , then linear decay to 0; jumping lr without warmup often spikes loss on step 1. MLM pre-training loss on masked positions only; fine-tune CE averages over batch labels, a much smaller gradient per token than MLM.

[CLS] token id is always position 0 after tokenization; pooling elsewhere grabs a punctuation embedding and tank accuracy silently. Max sequence length 128 truncates tail words in long reviews; for SST-2 most sentences fit, but log truncation rate on your split before trusting metrics. Whole-word masking groups WordPiece tokens; fine-tune does not use masking, but tokenizer still splits “running” into `run + ##ning`, affecting rare-word generalization. Pair-classification (MNLI) needs two segments with `token_type_ids` 0/1; SST-2 single segment uses all-zero type ids.

11 PyTorch Implementation Notes

- Hugging Face sequence classifier:

```
BertForSequenceClassification.from_pretrained(
    "bert-base-uncased", num_labels=2)
```

- Pass `attention_mask` (1 for real tokens, 0 for pad) and `token_type_ids` if model expects them.
- ```
AutoTokenizer(..., truncation=True, padding="max_length", max_length=128)
```

For batch encoding.

- Freeze lower layers optionally: `for p in model.bert.encoder.layer[:6].parameters(): p.requires_grad = False.`

- Use `get_linear_schedule_with_warmup` with optimizer step count matched to dataloader length  $\times$  epochs.

```
outputs = model(input_ids, attention_mask=attention_mask, labels=labels)
loss = outputs.loss # includes classification head
loss.backward()
torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
```

## 12 Local GPU Constraints

**bert-base** fine-tune at  $T = 128$ ,  $B = 8$  fits 4 GB with fp32; use  $B = 4$  if OOM during backward. **bert-large** is impractical full fine-tune on 1650; use **Module 11** LoRA or smaller distilled models.

Enable fp16 with `TrainingArguments(fp16=True)` in Hugging Face Trainer on Pascal+ GPUs. Gradient accumulation 4 with  $B = 4$  gives effective batch 16 without linear memory growth.

DistilBERT retains 97% of **bert-base** quality at 40% size; swap in when 4 GB blocks full BERT fine-tune before trying LoRA. Attention head probing: some heads specialize in delimiter tokens ([SEP], punctuation); pruning heads rarely helps SST-2 but matters in structural tasks.

## 13 Debugging Checklist

- All predictions class 0: wrong label mapping or learning rate too high destroying pretrained weights.
- Loss NaN: fp16 overflow; try full fp32 first epoch or lower lr.
- Worse than linear baseline: max length too short truncating sentiment cues, or forgetting attention mask.
- Slow training on 1650: dataloader bottleneck; cache tokenized tensors to disk.
- Checkpoint load error: mismatched `num_labels` or wrong tokenizer vocabulary file.

CLS logits random while MLM head works: classification head not trained long enough or lr  $10\times$  too high for head-only warmup phase.

## 14 Practice Problems

**Problem 1.** How many MLM labels per sequence of length 128 if 15% of tokens are masked (approx)?

**Problem 2.** Fine-tune with lr  $10^{-4}$  vs  $3 \times 10^{-5}$ ; plot val accuracy vs epoch for both.

**Problem 3.** Freeze bottom 6 layers; compare trainable param count and peak VRAM to full fine-tune.

**Problem 4.** Extract [CLS] embeddings for 500 sentences and visualize with PCA colored by label.

**Problem 5.** Compare last-hidden-state mean pooling vs [CLS] pooling on SST-2 without changing the pretrained backbone.

**Problem 6.** Tokenize the same sentence with and without `clean_up_tokenization_spaces=False`. Confirm [CLS] position and length stay stable.

Dynamic padding (`padding=True` without fixed `max_length`) saves compute but requires passing `attention_mask` every forward pass. Evaluate with `model.eval()` and disable dropout; forgotten `eval()` on BERT drops val accuracy 2 to 5 points on small sets. Export `save_pretrained`

after fine-tune includes config and tokenizer; loading with mismatched `num_labels` silently reshapes the classifier head. Run validation with identical `max_length` and truncation policy as training; eval-only longer context changes score without retraining. Distilled models (`distilbert-base-uncased`) trade 3 to 4 points accuracy for 40% faster inference on CPU baselines. Cache tokenized datasets to disk after first epoch; tokenization dominates wall time on SST-2 with naive Hugging Face loops. Report both accuracy and F1 on imbalanced dev sets; accuracy alone hides majority-class predictors.

Complete **Lab 09**; contrast with causal fine-tuning in **Module 10** and LoRA in **Module 11**.

## Source certification

This module lists where each core definition, equation, figure, and summary table comes from. Pedagogical simplifications (lab batch sizes, epoch counts, illustrative plots) are labeled in extension sections and are not claimed as optimal production settings.

**Table 3:** Certified topic map for Module 09.

| Topic                     | Primary source | Where in this PDF |
|---------------------------|----------------|-------------------|
| BERT architecture and MLM | [1]            | Eqs. 3–2          |
| Transformer encoder base  | [5]            | Section 2         |
| RoBERTa training notes    | [3]            | Extension section |

Full bibliographic entries appear below. Figures are schematic unless a caption cites a specific paper figure. If a statement is not listed here, treat it as general ML practice or lab-specific guidance, not a cited theorem.

## References

- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019.
- [2] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- [3] Yinhan Liu, Myle Ott, Naman Goyal, et al. RoBERTa: A robustly optimized BERT pretraining approach. *arXiv preprint arXiv:1907.11692*, 2019.
- [4] Alec Radford, Jeffrey Wu, Rewon Child, et al. Language models are unsupervised multitask learners. *OpenAI Technical Report*, 2019.
- [5] Ashish Vaswani, Noam Shazeer, Niki Parmar, et al. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.