

15. RL MDPs Q-Learning

Overview

Reinforcement learning formalizes decision making as Markov Decision Processes. We define returns, value functions, Bellman equations, and tabular Q-learning with ϵ -greedy exploration. This is the base for deep Q-networks in the next module.

Lab: Lab 15.

Contents

1	Sequential Decision Making	2
2	Markov Decision Process	2
3	Value Functions	2
4	Tabular Q-Learning	2
5	Exploration vs. Exploitation	2
6	Policy and Value Iteration	3
7	Lab	3
8	SARSA and On-Policy Control	3
8.1	Eligibility Traces	3
9	Extended Intuition	3
10	Numerical Walkthrough	4
11	PyTorch Implementation Notes	4
12	Local GPU Constraints	4
13	Debugging Checklist	4
13.1	MDP Design Choices Engineers Actually Control	5
14	Practice Problems	5

Module track

This module starts a new track. Later modules refer back using **Module NN** labels.



Figure 1: Module sequence (solid box: this PDF; dashed: typical next step).

1 Sequential Decision Making

Reinforcement learning [2] optimizes cumulative reward in environments with state transitions. Differs from supervised pipelines in **Module 09**.

2 Markov Decision Process

MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, P, R, \gamma)$:

$$P(s_{t+1} \mid s_t, a_t), \quad r_t = R(s_t, a_t), \quad \gamma \in [0, 1]. \quad (1)$$

Markov property: $P(s_{t+1} \mid s_t, a_t, s_{t-1}, \dots) = P(s_{t+1} \mid s_t, a_t)$.

Note. γ discounts future rewards; $\gamma \rightarrow 0$ myopic, $\gamma \rightarrow 1$ far-sighted (requires episodic termination).

3 Value Functions

State-value and action-value:

$$V^\pi(s) = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t = s \right], \quad (2)$$

$$Q^\pi(s, a) = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k r_{t+k} \mid s_t = s, a_t = a \right]. \quad (3)$$

Optimal Q^* satisfies Bellman optimality:

$$Q^*(s, a) = \mathbb{E} \left[R + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]. \quad (4)$$

4 Tabular Q-Learning

Model-free update for discrete $\mathcal{S}, \mathcal{A}$:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]. \quad (5)$$


Off-policy: behavior may be ϵ -greedy while learning greedy target.

5 Exploration vs. Exploitation

ϵ -greedy: choose random action with prob. ϵ , else $\arg \max_a Q(s, a)$. Decay ϵ from 1.0 to 0.01 over training.

Table 1: RL algorithm taxonomy.

Method	Policy	Value
Q-learning	Implicit (greedy)	Tabular Q
DQN (Module 16 (DQN loss))	Implicit	Neural Q
Policy gradient (Module 17 (policy gradient theorem))	Explicit π_θ	Optional

**Figure 2:** Agent-environment interaction loop.

6 Policy and Value Iteration

Dynamic programming backups when P is known:

$$V_{k+1}(s) = \max_a \sum_{s'} P(s'|s, a) [R(s, a) + \gamma V_k(s')]. \quad (6)$$

7 Lab

Lab 15 implements tabular Q-learning on GridWorld using (5). Neural extension in **Module 16** (DQN loss).

8 SARSA and On-Policy Control

SARSA update uses actual next action a' from behavior policy:

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma Q(s', a') - Q(s, a)]. \quad (7)$$

Safer under exploration than Q-learning when ϵ -greedy explores dangerous states.

8.1 Eligibility Traces

TD(λ) blends n -step returns: $G_t^{(\lambda)} = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$.

9 Extended Intuition

Reinforcement learning models sequential decision making: an agent observes state s_t , takes action a_t , receives reward r_{t+1} , transitions to s_{t+1} . Markov Decision Processes assume the future depends on the present state only, not the full history (if state is sufficient). The return $G_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ discounts future rewards with $\gamma \in [0, 1)$.

Value functions estimate expected return: $V^\pi(s)$ under policy π , or action-value $Q^\pi(s, a)$. Optimal Q^* satisfies the Bellman optimality equation:

$$Q^*(s, a) = \mathbb{E} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]. \quad (8)$$

Tabular Q-learning updates with learning rate α :

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]. \quad (9)$$

ϵ -greedy exploration trades random actions (explore) with greedy actions (exploit).

Deep methods in **Module 16** and **Module 17** scale this to neural function approximators; start tabular to verify the math.

10 Numerical Walkthrough

GridWorld 4×4, $\gamma = 0.9$, $\alpha = 0.5$, $\epsilon = 0.1$, goal at bottom-right reward +1, step cost -0.04 .

Start $Q(s, a) = 0$ everywhere. From state s (row 0, col 0), action right, land in s' with $r = -0.04$:

$$Q(s, \text{right}) \leftarrow 0 + 0.5 \left(-0.04 + 0.9 \max_{a'} Q(s', a') - 0 \right) = 0.5(-0.04) = -0.02. \quad (10)$$

After many episodes, Q near goal states rises toward $1 - 0.04/(1 - \gamma)$ scaled by path length.

Episode length maybe 20 to 50 steps early, dropping toward optimal 6 steps as policy improves. With $\epsilon = 0.1$, roughly 10% random moves persist even after convergence (continued exploration).

1000 episodes on 4×4: tabular update $\approx$ 15k steps total; runs in milliseconds on CPU.

11 PyTorch Implementation Notes

Tabular lab uses NumPy dict or `np.zeros((n_states, n_actions))`; no GPU needed.

- State indexing: flatten (*row, col*) to single index for table lookup.
- `np.argmax(Q[s])` for greedy action; `random` branch with prob ϵ .
- Log episode return `sum(rewards)` and steps per episode for learning curves.
- For later DQN (**Module 16**): same env API (`gymnasium`) eases transition.
- Reproducibility: `env.reset(seed=...)` and `np.random.seed`.

```
q = np.zeros((n_states, n_actions))
td_target = r + gamma * q[s_next].max()
q[s, a] += alpha * (td_target - q[s, a])
```

12 Local GPU Constraints

Tabular Q-learning is CPU-only; no VRAM concerns. If prototyping neural Q on the same GridWorld in **Module 16**, tiny MLP fits any GPU.

Use vectorized env wrappers only when moving to DQN at scale; tabular code stays readable with explicit loops.

13 Debugging Checklist

- Q values diverge to large magnitude: α too high or missing terminal state bootstrap ($\max Q = 0$ at goal).
- Policy never reaches goal: ϵ too low early or rewards too sparse; increase exploration or shape rewards.

- Oscillating policy: non-stationary if updating while acting on stale Q without proper episode boundaries.
- Same action always: tie-breaking not random in `argmax` on equal Q values.
- Off-by-one in s' after action: coordinate system does not match env `step` return.

13.1 MDP Design Choices Engineers Actually Control

The MDP is not handed to you by the environment API; you define what enters the state vector, how rewards are scaled, and when episodes terminate. A sparse goal reward (+1 at target) teaches slowly because most transitions look identical until the agent stumbles on success. Reward shaping adds dense signals (distance-to-goal penalty) but can change the optimal policy if shaping rewards are not potential-based; log both shaped and sparse returns when experimenting. Continuing tasks (robot locomotion) use undiscounted or γ very close to 1; episodic tasks (GridWorld, games) terminate and bootstrap $Q(s', a') = 0$ at terminal states.

Partial observability appears when the agent sees less than the full Markov state (frame pixels without velocity). Wrapping history into the state (stack last k observations) restores Markov property at the cost of larger tables or networks later in **Module 16**. Action space design matters: factorizing movement into 4 directions keeps the tabular grid tractable; continuous torque control needs policy methods from **Module 17**.

Policy evaluation computes expected return under a fixed policy; policy improvement takes $\arg \max_a Q(s, a)$ and iterates toward π^* . Q-learning is off-policy: it learns about the greedy policy while following ϵ -greedy behavior, which is why exploration schedule remains important even after Q looks converged. Discount $\gamma = 0.99$ vs 0.9 changes how many steps ahead the agent plans; on GridWorld, lower γ favors shorter paths even if per-step penalty is small. Always log episode length, total return, and visit counts per state when debugging; a policy that oscillates between two cells often signals incorrect terminal handling or reward sign errors. Seed environments and RNG for reproducible learning curves before comparing α or ϵ schedules across runs. Windy GridWorld variants (random action with prob 0.2 regardless of policy) teach robustness testing: optimal tabular policy should still reach goal in expectation. Sutton example 6.6 style: if step cost is -1 and goal reward is 0, optimal policy minimizes steps; sign errors in reward definition flip the learned path completely. Monte Carlo methods wait until episode end to update; TD methods bootstrap after one step. Tabular Q-learning is TD(0); compare update frequency when debugging slow convergence. State aggregation (grouping 4×4 grid into 2×2 macro cells) reduces table size but loses Markov property if critical detail sits inside a macro cell. Frozen policy evaluation: after Q converges, sweep $\epsilon = 0$ for 100 episodes and report mean return; compare to random policy baseline to quantify improvement. Stochastic transitions (slippery grid with 0.8 intended direction, 0.1 each sideways) require averaging over next states in the Bellman backup; deterministic env code will under-estimate risk. Action masking for invalid moves (walls) should zero out Q updates for impossible actions rather than letting agent waste exploration on illegal steps. SARSA (on-policy TD) updates with $Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma Q(s', a') - Q(s, a)]$ where a' is the action actually taken; compare to Q-learning when environment is highly stochastic.

14 Practice Problems

Problem 1. With $\gamma = 0.95$, step penalty -1 , goal reward $+10$ at 5 steps optimal path, approximate optimal $Q(s_0, a_0)$.

Problem 2. Implement ϵ decay from 1.0 to 0.05 over 500 episodes; plot return vs episode.

Problem 3. Compare synchronous updates vs updating Q while stepping within the same episode.

Problem 4. Visualize greedy policy arrows on the grid after training.

Implement in **Lab 15**; neural scaling in **Module 16**; policy gradients in **Module 17**.

Source certification

This module lists where each core definition, equation, figure, and summary table comes from. Pedagogical simplifications (lab batch sizes, epoch counts, illustrative plots) are labeled in extension sections and are not claimed as optimal production settings.

Table 3: Certified topic map for Module 15.

Topic	Primary source	Where in this PDF
MDP / Bellman equations	[2]	Eqs. 4 to 4
Tabular Q-learning	[3]	Section 4

Full bibliographic entries appear below. Figures are schematic unless a caption cites a specific paper figure. If a statement is not listed here, treat it as general ML practice or lab-specific guidance, not a cited theorem.

References

- [1] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, et al. Human-level control through deep reinforcement learning. *Nature*, 518:529–533, 2015.
- [2] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018.
- [3] Christopher J. C. H. Watkins and Peter Dayan. Q-learning. *Machine Learning*, 8:279–292, 1992.
- [4] Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. volume 8, pages 229–256, 1992.