

16. DQN

Overview

Deep Q-Networks approximate Q-values with neural nets, using experience replay and target networks for stability. We derive the TD loss, discuss overestimation bias, and map hyperparameters to CartPole training on CPU/GPU.

Lab: Lab 16.

Prior modules: Module 15.

Contents

1	Scaling Q-Learning	2
2	DQN Architecture	2
3	DQN Loss	2
4	Experience Replay	2
5	ϵ-Greedy and Double DQN	2
6	Limitations	2
7	Lab	3
8	Prioritized Replay	3
8.1	Rainbow Combine	3
9	Extended Intuition	3
10	Numerical Walkthrough	4
11	PyTorch Implementation Notes	4
12	Local GPU Constraints	4
13	Debugging Checklist	4
13.1	Replay Buffer Mechanics and DQN Variants	5
14	Practice Problems	5

Module track

This module builds on **Module 15**. Read those PDFs first. References to other modules appear as **Module NN** (not page numbers, because each module is its own PDF).




Figure 1: Module sequence (solid box: this PDF; dashed: typical next step).

1 Scaling Q-Learning

Tabular **Module 15** (equation) fails on high-dimensional states (images, Atari frames). DQN [3] approximates $Q(s, a; \theta)$ with a CNN.

2 DQN Architecture

Input stack of 4 grayscale frames $\rightarrow$ conv layers $\rightarrow$ fully connected $\rightarrow |\mathcal{A}|$ Q-values.

$$Q(s, a; \theta) \approx Q^*(s, a). \quad (1)$$

Two key innovations. Experience replay and target network stabilize bootstrapping from **Module 15** (Bellman equation).

3 DQN Loss

Sample minibatch (s, a, r, s') from replay buffer $\mathcal{D}$:

$$\mathcal{L}(\theta) = \mathbb{E}_{(s, a, r, s') \sim \mathcal{D}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2 \right]. \quad (2)$$

θ^- is target network updated every C steps: $\theta^- \leftarrow \theta$.

4 Experience Replay

Store transitions in circular buffer capacity 10^6 ; uniform random sampling breaks temporal correlation. Improves sample efficiency over on-line updates from **Module 15** (Q-learning update).

5 ϵ -Greedy and Double DQN

Double DQN decouples selection and evaluation:

$$y = r + \gamma Q\left(s', \arg \max_{a'} Q(s', a'; \theta); \theta^-\right). \quad (3)$$


Reduces overestimation bias from max operator in (2).

6 Limitations

DQN handles discrete actions only; continuous control needs policy gradients (**Module 17** (policy gradient theorem)). Reward clipping and frame preprocessing are domain-specific engineering.

Table 1: DQN hyperparameters (Atari).

Hyperparameter	Value
Replay buffer size	10^6
Minibatch size	32
γ	0.99
Target update freq. C	10,000 steps
ϵ anneal	$1.0 \rightarrow 0.1$

**Figure 2:** DQN with experience replay.

7 Lab

Lab 16 trains DQN on CartPole or Atari; compare convergence to tabular Q in **Lab 15**.

8 Prioritized Replay

Sample transition i with probability $P(i) \propto |\delta_i| + \epsilon$:

$$w_i = \left(\frac{1}{N} \cdot \frac{1}{P(i)} \right)^\beta, \quad \text{importance-weighted loss.} \quad (4)$$

Dueling architecture separates $V(s)$ and advantage $A(s, a)$: $Q(s, a) = V(s) + A(s, a) - \mathbb{E}_a A(s, a)$.

8.1 Rainbow Combine

Prioritized + Double + Dueling + Multi-step + Distributional + Noisy nets stack improvements over vanilla (2).

9 Extended Intuition

Deep Q-Networks replace the Q-table with $Q_\theta(s, a)$ parameterized by a neural network. Experience replay stores transitions (s, a, r, s') in a buffer and samples random mini-batches, breaking temporal correlation that destabilizes SGD. Target network Q_{θ^-} uses frozen weights updated periodically so bootstrap targets do not chase a moving objective every step.

Loss for one sample:

$$\mathcal{L} = \left(r + \gamma \max_{a'} Q_{\theta^-}(s', a') - Q_\theta(s, a) \right)^2. \quad (5)$$

Double DQN reduces overestimation by selecting a' with online net, evaluating with target net (optional extension).

Atari adds frame stacking, reward clipping, and convolutional encoders; CartPole uses a 2-layer MLP on 4-dim state. Exploration remains ϵ -greedy, often annealed over 100k to 1M frames.

Policy gradient methods in **Module 17** handle continuous action spaces where $\max_{a'}$ is impractical.

10 Numerical Walkthrough

CartPole-v1: state dim 4, 2 discrete actions, MLP $4 \rightarrow 128 \rightarrow 128 \rightarrow 2$.

Replay buffer capacity 100k transitions; batch $B = 64$, $\gamma = 0.99$, lr = 10^{-3} , target update every 1000 steps.

Single transition: s , action left, $r = 1.0$, s' not done. If $Q_{\theta-}(s', \cdot) = [10.2, 9.8]$, bootstrap = $0.99 \cdot 10.2 = 10.098$. If $Q_{\theta}(s, \text{left}) = 8.5$, squared error $(1 + 10.098 - 8.5)^2 \approx 6.8$.

Training 500 episodes: moving average return crosses 200 (solved threshold) often by episode 150 to 300 with tuned ϵ schedule. Frame skip not used in CartPole; Atari would skip 4 frames and max-pool.

Memory: buffer $100k \times (4 + 1 + 1 + 4 + 1)$ floats negligible; network params $\approx 50k$.

11 PyTorch Implementation Notes

- Create env with `gymnasium.make("CartPole-v1")`.
- Step returns `obs, reward, terminated, truncated, info = env.step(a)`.
- Store `done = terminated or truncated` for bootstrap mask.
- Q-network outputs vector of Q-values per action; pick action ϵ -greedy on `q_values.argmax()`.
- Target computation:

```
with torch.no_grad():
    next_q = target_net(s_next).max(1)[0]
    target = r + gamma * next_q * (1 - done)
```

- Gather predicted Q: `q_sa = q_values.gather(1, a.unsqueeze(1)).squeeze(1)`.

```
loss = F.mse_loss(q_sa, target)
optimizer.zero_grad()
loss.backward()
optimizer.step()
if step % 1000 == 0:
    target_net.load_state_dict(q_net.state_dict())
```

12 Local GPU Constraints

CartPole DQN trains on CPU in seconds; GPU optional for habit. Atari CNN DQN on 1650: use frame size 84x84, replay on CPU RAM, batch 32; may need hours for 1M frames.

Mixed precision rarely needed for small MLP DQN; Atari benefits modestly from GPU conv forward.

13 Debugging Checklist

- Return stuck below 50: ϵ not decaying, or target net never updates.
- Q-values explode: lr too high; clip gradients or reduce lr to 10^{-4} .
- Loss zero immediately: replay buffer empty or same transition repeated; verify `buffer.sample`.
- Done flag ignored: bootstrap at terminal states inflates targets.
- Action always same: network outputs tied; check initialization and exploration.

13.1 Replay Buffer Mechanics and DQN Variants

The replay buffer is a FIFO queue: when capacity is exceeded, oldest transitions drop off, which slowly removes bias toward early poor policies. Uniform sampling assumes all transitions are equally informative; prioritized replay (optional extension) weights samples by TD error magnitude so surprising transitions get replayed more often. Store transitions as $(\mathbf{s}, \mathbf{a}, \mathbf{r}, \mathbf{s}', \text{done})$ tuples in float32; for Atari, stack 4 frames into s before pushing to keep Markov structure without storing full video.

Target network sync every C steps (1000 to 10000) trades stability for lag: too frequent sync chases a moving target; too rare sync slows learning. Double DQN decouples selection and evaluation:

$$y = r + \gamma Q_{\theta^-} \left(s', \arg \max_{a'} Q_{\theta}(s', a') \right). \quad (6)$$

This reduces systematic overestimation from max operator bias in vanilla DQN. Dueling architecture splits $Q(s, a) = V(s) + A(s, a) - \text{mean}_a A(s, a)$, helping states where action choice matters little (many Atari frames).

ϵ -greedy schedule: start $\epsilon = 1.0$, linear decay to 0.05 over first 50k to 100k steps for CartPole; hold small ϵ forever to avoid premature exploitation on stochastic envs. Frame stacking for Atari: each s_t concatenates grayscale frames $t - 3, \dots, t$ along channel dim; apply max over last two raw frames on s_t when a_t is fire button to reduce flicker. Reward clipping to $\{-1, 0, +1\}$ in Atari stabilizes Q magnitudes but removes fine-grained feedback; skip clipping on CartPole where rewards are already bounded. Gradient clipping at norm 10 before optimizer step prevents occasional TD target spikes from destabilizing weights when replay contains rare high-return trajectories. Compare final sample efficiency against tabular Q from **Module 15** on the same discretized CartPole grid to verify the network actually generalizes across similar states. n -step returns ($n = 3$) blend one-step TD and Monte Carlo: target uses reward sequence $r_{t+1}, \dots, r_{t+n}$ plus bootstrap $Q(s_{t+n})$; reduces bias at cost of slightly higher variance. Huber loss on TD error dampens impact of outlier targets compared to pure MSE; switch when Q-values occasionally spike above 100 on Atari runs. Soft target updates (Polyak averaging $\theta^- \leftarrow \tau\theta + (1 - \tau)\theta^-$) replace hard sync every C steps in some implementations; $\tau = 0.005$ is a common default. Set `torch.backends.cudnn.benchmark = True` for fixed-size Atari conv stacks; disable when input shapes vary per batch to avoid autotuner overhead. Minimum replay size before training starts: wait until buffer holds 1000+ transitions or Q overfits to correlated early episodes. Noisy nets (optional) add learned parametric noise to weights for exploration without ϵ -greedy; useful when action space is large and epsilon schedule is hard to tune. Log TD error histogram monthly in long Atari runs; bimodal distribution (many near zero, tail of spikes) suggests prioritized replay would help. Save Q-network weights every 50k frames; Atari runs crash from driver timeouts and replay buffer rebuild is expensive.

14 Practice Problems

Problem 1. Replay buffer size 50k, sample batch 64; expected unique transitions per batch (approx)?

Problem 2. Plot Q-value of best action at initial CartPole state over training steps.

Problem 3. Ablate target network (update every step); compare stability to periodic sync.

Problem 4. Compare tabular Q on discretized CartPole (**Module 15**) vs DQN sample efficiency.

Train in **Lab 16**; builds on **Module 15**; compare to policy gradients in **Module 17**.

Source certification

This module lists where each core definition, equation, figure, and summary table comes from. Pedagogical simplifications (lab batch sizes, epoch counts, illustrative plots) are labeled in extension sections and are not claimed as optimal production settings.

Table 3: Certified topic map for Module 16.

Topic	Primary source	Where in this PDF
DQN TD target	[3]	Eq. 3
Experience replay	[3]	Section 4
Double DQN bias note	[5]	Extension section
RL foundations	[4]	Overview

Full bibliographic entries appear below. Figures are schematic unless a caption cites a specific paper figure. If a statement is not listed here, treat it as general ML practice or lab-specific guidance, not a cited theorem.

References

- [1] Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- [2] Volodymyr Mnih, Adria Puigdomenech Badia, Mehdi Mirza, et al. Asynchronous methods for deep reinforcement learning. In *ICML*, 2016.
- [3] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, et al. Human-level control through deep reinforcement learning. *Nature*, 518:529–533, 2015.
- [4] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2nd edition, 2018.
- [5] Hado Van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double Q-learning. In *AAAI*, 2016.