

20. SHAP LIME XAI Deployment

Overview

Explainability methods attribute predictions to input features. SHAP uses Shapley values from cooperative game theory; LIME fits local interpretable surrogates. The module closes with serving considerations: logging explanations alongside predictions in a small Flask deployment.

Lab: Lab 20.

Prior modules: Module 11.

Contents

1	Why Explainability?	2
2	SHAP Values	2
3	Shapley Kernel	2
4	LIME	2
5	Model-Specific Accelerations	2
6	Deployment Integration	2
7	Limitations	3
8	Lab	3
9	Global vs. Local Explanations	3
9.1	Regulatory Context	3
10	Extended Intuition	3
11	Numerical Walkthrough	4
12	PyTorch Implementation Notes	4
13	Local GPU Constraints	5
14	Debugging Checklist	5
14.1	Explanation Fidelity, Correlation, and Monitoring	5
15	Practice Problems	6

Module track

This module builds on **Module 11**. Read those PDFs first. References to other modules appear as **Module NN** (not page numbers, because each module is its own PDF).




Figure 1: Module sequence (solid box: this PDF; dashed: typical next step).

1 Why Explainability?

Deployed models from **Module 19** (MLflow tracking) affect users; regulators and stakeholders demand interpretability. Post-hoc methods explain black-box predictions without changing model internals.

2 SHAP Values

SHAP [3] assigns feature attributions ϕ_i satisfying efficiency:

$$\sum_{i=1}^M \phi_i = f(\mathbf{x}) - \mathbb{E}[f(\mathbf{X})], \quad (1)$$

derived from Shapley values in cooperative game theory.

Note. SHAP provides unified framework: LIME (section 4), DeepLIFT, and TreeSHAP are approximations under specific assumptions.

3 Shapley Kernel

Weighted linear regression on coalition indicators $\mathbf{z} \in \{0, 1\}^M$:

$$\pi_{\mathbf{x}}(\mathbf{z}) = \frac{(M-1)!}{\binom{M}{|\mathbf{z}|} |\mathbf{z}|! (M-|\mathbf{z}|)!}. \quad (2)$$

Kernel SHAP estimates (1) by sampling coalitions.

4 LIME

LIME [5] fits local interpretable model g near $\mathbf{x}$:

$$\xi(\mathbf{x}) = \arg \min_{g \in \mathcal{G}} \mathcal{L}(f, g, \pi_{\mathbf{x}}) + \Omega(g), \quad (3)$$

where $\mathcal{L}$ measures fidelity of g to black-box f on perturbed samples weighted by $\pi_{\mathbf{x}}$.

5 Model-Specific Accelerations

TreeSHAP: exact Shapley for tree ensembles in polynomial time. Gradient SHAP: integrates gradients for neural nets fine-tuned via LoRA in **Module 11** (LoRA definition).

6 Deployment Integration

Serving pipeline returns $(\hat{y}, \phi)$ for audit logs. Monitor attribution drift when input distribution shifts (see **Module 19** (MLflow tracking)).


Table 1: XAI method selection.

Model	Method	Scope
Any classifier	LIME	Local
Tree ensembles	TreeSHAP	Global + local
Deep networks	Gradient SHAP	Local
LLM	Attention rollout	Token-level

7 Limitations

Explanations are not causal; correlated features split attribution arbitrarily. Adversarial perturbations can fool LIME/SHAP without changing prediction.

Consistency: if $f(\mathbf{x}) = f(\mathbf{x}')$, SHAP assigns same ϕ_i when contributions identical. (4)

**Figure 2:** Waterfall-style SHAP contributions to $f(\mathbf{x})$.

8 Lab

Lab 20 computes SHAP and LIME on a deployed classifier; compare attributions from (1) and (3). Integrate explanation logging into the serving stack from **Lab 19**.

9 Global vs. Local Explanations

SHAP summary plots aggregate $|\phi_i|$ across dataset; partial dependence shows marginal effect:

$$\text{PD}_j(x_j) = \mathbb{E}_{\mathbf{x}_{\setminus j}}[f(x_j, \mathbf{x}_{\setminus j})]. \quad (5)$$

Counterfactual explanations: minimal change δ s.t. $f(\mathbf{x} + \delta) \neq f(\mathbf{x})$.

9.1 Regulatory Context

EU AI Act high-risk systems may require documentation of explanation methodology and failure modes.

10 Extended Intuition

SHAP attributes a prediction to input features using Shapley values from cooperative game theory: fair distribution of payout across players (features). For tree models, TreeSHAP is fast and exact;

for neural nets, KernelSHAP or DeepSHAP approximate with background samples. LIME fits a local linear surrogate around one prediction by perturbing inputs and weighting samples by proximity.

Global importance (mean $|\text{SHAP}|$) differs from local explanations for a single row; both answer different stakeholder questions. Explanation methods can disagree, especially under correlated features; never treat attributions as causal.

Deployment adds latency and storage: log top- k features per request, cache explanations for batch offline review, or compute async. Fine-tuned models from **Module 11** and training tracked in **Module 19** should expose the same preprocessing in serving and explanation paths.

Regulatory contexts may require human-readable reason codes derived from ranked attributions.

11 Numerical Walkthrough

Binary classifier on 5 tabular features; one instance $x = [1.2, 0.0, 3.4, -0.5, 2.1]$, prediction $p = 0.82$. SHAP values $\phi = [0.15, -0.05, 0.30, 0.02, 0.10]$ sum with base value $\phi_0 = 0.30$ to logit or prob link: $0.30 + 0.15 - 0.05 + 0.30 + 0.02 + 0.10 = 0.82$ (prob space illustration; exact link depends on model output).

LIME: generate 500 perturbed samples around x (Gaussian on numeric features), get predictions p_i , fit weighted linear model with weights $K(x, x_i) = \exp(-\|x - x_i\|^2/\sigma^2)$, $\sigma = 0.5$. Surrogate $R^2 = 0.75$ on neighborhood indicates local linear approximation quality.

Latency budget: TreeSHAP 2 ms, KernelSHAP 200 ms for same model; batch 100 explanations for offline audit vs real-time API.

Flask endpoint returns `{"prediction": 1, "shap": {"f0": 0.15, ...}}` JSON; log alongside input hash for GDPR-aware retention policies.

12 PyTorch Implementation Notes

- `shap.Explainer(model, background_data)` auto-selects algorithm; for sklearn trees use `TreeExplainer`.
- Deep explainer: background tensor shape (50, C, H, W) for images; keep background small on 1650.
- `shap_values = explainer(x)`; force plot with `shap.plots.waterfall` for one row.
- LIME: `lime.lime_tabular.LimeTabularExplainer(training_data, feature_names=..., class_names=...)`.
- Wrap PyTorch model with `predict_fn` that applies softmax and returns numpy for SHAP/LIME.

```
def predict_fn(x):
    model.eval()
    with torch.no_grad():
        t = torch.tensor(x, dtype=torch.float32, device=device)
        return model(t).softmax(1).cpu().numpy()

explainer = shap.KernelExplainer(predict_fn, background_numpy)
sv = explainer.shap_values(instance, nsamples=100)
```

13 Local GPU Constraints

KernelSHAP with 1000 coalitions on large neural nets is slow on 1650; reduce `nsamples` to 50 to 100 for demos. Background dataset 50 to 100 rows balances fidelity vs time; use k-means summary of train set (`shap.kmeans`).

Image DeepSHAP at 224px: batch background 8, expect seconds per explanation; run offline not in request path. LIME on tabular is CPU-cheap; prefer for real-time alongside tree models.

14 Debugging Checklist

- SHAP sum does not match prediction: wrong link function (logit vs prob) or missing base value.
- All zero attributions: model constant on perturbations; check preprocessing mismatch vs training.
- LIME unstable across runs: increase samples or fix random seed on perturbations.
- Explainer OOM on GPU: move background and explain batch to CPU for KernelSHAP.
- Serving mismatch: training scaler not applied in `predict_fn` used by explainer.

14.1 Explanation Fidelity, Correlation, and Monitoring

SHAP values satisfy efficiency: $\sum_i \phi_i + \phi_0 = f(x)$ for the chosen model output link (logit or probability). TreeSHAP exploits tree structure for exact Shapley values in polynomial time; KernelSHAP treats the model as a black box and samples coalitions, trading accuracy for generality. Background dataset defines the reference expectation ϕ_0 ; a background of all zeros on standardized features is wrong if training mean was non-zero. Use `shap.kmeans(train, 50)` to summarize background when raw train has 100k rows; 50 centroids usually capture enough distribution for KernelSHAP.

LIME trains weighted linear surrogate g locally: minimize $\sum_i \pi_x(z_i)(f(z_i) - g(z_i))^2 + \Omega(g)$ where π_x is proximity kernel. High R^2 on perturbed samples means local linearity holds; low R^2 warns the explanation may mislead for that instance. Correlated features (zip code and city) receive split credit that is not uniquely identifiable; report grouped explanations or drop one feature from the set when presenting to stakeholders.

Production patterns: async explanation queue for batch review, sync TreeSHAP for low-latency tree models, precompute SHAP for static batch scoring jobs. Log explanation payloads with model version from MLflow (**Module 19**) so drift investigations tie back to the exact checkpoint. Monitor mean $|\phi_i|$ per feature over sliding windows; sudden spikes in a feature's average attribution often precede data drift alerts. Adversarial inputs can produce misleading attributions; cross-check with holdout performance and LIME on the same row before trusting a single method. For neural models fine-tuned with LoRA (**Module 11**), run explainer on merged weights so `predict_fn` matches serving graph exactly. GDPR-style retention: store hashed inputs and top-3 attributions, not full SHAP vectors, when regulations limit sensitive feature logging.

Interventional vs observational SHAP: interventional breaks correlations by replacing feature values from background independently; observational preserves joint distribution. Pick interventional when stakeholders need "what if we change this one input"; pick observational when features naturally co-occur and you want realistic counterfactuals. Force plots and waterfall charts communicate local attributions to non-technical reviewers better than raw ϕ vectors. Batch explanation jobs should cap `nsamples` per row and shard across workers; log wall time per 1000 rows for capacity planning. SHAP interaction values $\Phi_{i,j}$ capture pairwise feature effects; compute only for top features when $p > 20$ to avoid exponential coalition cost. LIME for text: perturb by

deleting word subsets, not Gaussian noise; use `LimeTextExplainer` with same tokenizer as model input pipeline. Fairness review: compare mean SHAP of sensitive attribute across demographic slices; large disparity triggers model audit even when AUC looks acceptable. Explanation cache key: hash (model_version, preprocessing_hash, input_row); invalidate cache on any scaler or vocab change from **Module 19** deployment.

15 Practice Problems

Problem 1. For 3 features with given SHAP values, verify they sum to prediction minus base rate in logit space.

Problem 2. Compare SHAP vs LIME top-3 features on 10 test rows; measure agreement rate.

Problem 3. Add explanation logging to Flask app from **Lab 19** stack; measure p95 latency overhead.

Problem 4. Identify two correlated features where SHAP splits credit arbitrarily; discuss mitigation (grouped SHAP or drop one).

Complete **Lab 20**; integrates models from **Module 11** and tracking from **Module 19**.

Source certification

This module lists where each core definition, equation, figure, and summary table comes from. Pedagogical simplifications (lab batch sizes, epoch counts, illustrative plots) are labeled in extension sections and are not claimed as optimal production settings.

Table 3: Certified topic map for Module 20.

Topic	Primary source	Where in this PDF
SHAP Shapley values	[3]	Eq. 1
LIME local surrogate	[5]	Eq. 3
XAI survey context	[2]	Overview
Interpretability textbook notes	[4]	Extension section

Full bibliographic entries appear below. Figures are schematic unless a caption cites a specific paper figure. If a statement is not listed here, treat it as general ML practice or lab-specific guidance, not a cited theorem.

References

- [1] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. BERT: Pre-training of deep bidirectional transformers for language understanding. In *NAACL*, 2019.
- [2] Riccardo Guidotti, Anna Monreale, Salvatore Ruggieri, et al. A survey of methods for explaining black box models. *ACM Computing Surveys*, 2018.
- [3] Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *NeurIPS*, 2017.
- [4] Christoph Molnar. *Interpretable Machine Learning*. Leanpub, 2020.
- [5] Marco Tulio Ribeiro, Sameer Singh, and Carlos Guestrin. “why should i trust you?” explaining the predictions of any classifier. In *KDD*, 2016.